Instituto Superior de Engenharia do Porto
Departamento de Engenharia Informática

MSc in Informatics Engineering
Armazéns de Dados (ARMDD)
2026/2027 – 1st semester

MoreRhythm Data Mart

Practical Work – Part 1
Architecture, Dimensional Model, Staging Area, Mapping and Data Cleansing

André Gonçalves – 1230794
Michal Binek – 1260224

Porto, November 2026

Contents

1. Introduction

MoreRhythm, SA is a national company, headquartered in Porto, that distributes and sells musical instruments across Portugal and occasionally in Spain. Its operational system supports billing, purchasing and stock management. Management wants a **data mart for historical analysis of the sales made to customers**, at the lowest level of detail, so that any analysis can be built on it, including the 10 analyses listed in the project brief.

The brief points out three characteristics of the operational data that the solution has to handle:

- the same customer may have been registered more than once, sometimes with slightly different data;
- shipping costs are recorded per invoice and not per invoice line;
- every table records when each row was inserted and when it was last changed.

The data mart must also use the **euro as the standard currency**, because the company plans to expand to other EU countries and the United Kingdom. And the Integration Services project must have **no absolute paths**: servers, database names and paths come from parameters.

This report covers **Part 1** of the work: data profiling of the source, the data warehouse architecture, the dimensional model (Kimball methodology), the Staging Area structures, the mapping between source, Staging Area and Data Mart with the SCD strategy, and the data cleansing and transformation operations. The SQL scripts, diagrams and documentation are in the GitHub repository `ARMDD_26-27_1260224_1230794`.

The work was developed with SQL Server 2025, SQL Server Management Studio and Visual Studio 2026 (Integration Services), following the course methodology (Kimball dimensional modelling) and the conventions of the laboratory classes.

2. Data Profiling

Profiling was run in SSMS against the operational database `MoreRhythm` (script `01_Data_Profiling.sql`; raw output in `01_Data_Profiling_results.txt`).

2.1 Volumes

Table	Rows		Table	Rows
customers	26		products	67
invoices	37		suplliers	14
invoices_lines	93		employees	6

families	17		categories	11
zones	5		stocks_movements / purchases	686 / 18

The sample is small but, as stated in the brief, representative; the ETL is designed for large volumes (set-based SQL, indexed staging keys, incremental extraction).

2.2 Customer tax numbers (`customers.taxpayer_number`)

Finding	Records	Treatment
Placeholder tax numbers: `''` , `0` , `00` , `000000000`	20, 24, 25 (internal branches), 2 ("Vendas a Dinheiro")	Not an error: set `TaxNumber = NULL` , keep the customer; excluded from NIF matching
Exact duplicate NIF `502323434`	10 and 18 ("Eira e Beira, lda")	De-duplicate (rule M1)
Near-duplicate – NIF differs in 2 transposed digits, same zip, same/similar name	5 26 ("Coro Nossa Senhora Assunção" / "Coro N. Senhora Assunção"), 8 22 ("Luz e Som, lda" / "Luz e Som Lda.")	De-duplicate (rule M2)
NIF fails the Portuguese modulo-11 check digit	17 of 23 real NIFs	**Not** used as a rejection rule (would discard most of the sample); used only as a survivor tie-breaker and stored as `IsTaxNumberValid`

2.3 Other customer attributes

- 1 invalid e-mail (`coro.sapo.pt` , customer 5) set to NULL (the duplicate 26 has the valid `coro@sapo.pt` , recovered by survivorship).
- 12 customers without e-mail, 3 without city, 2 without zone (internal / cash-sale customers).
- Customer 21 (Spain): city empty, zip column holds `28108 Alcobendas - Madrid, España` city/zip parsed; country derived from zone 5 = "Espanha".
- Zip codes hold *code* + *locality* (`2600-011 VILA FRANCA DE XIRA`) split into `ZipCode` and `ZipLocality` .
- `seller` = 0 on 4 customers does not match any employee kept as attribute only (not a FK).

2.4 Invoice timestamps

- `invoice_date` range 2025-01-31 … 2025-05-01; invoices are issued at month end (31-Jan, 28-Feb, 31-Mar, 30-Apr) plus one on 1-May.
- `created_at ≤ updated_at` holds for all invoices and lines.
- 36 of 37 invoices were updated after insertion and **all** have `last_update_date = payment_date` invoices change after the sale (payment registration). The incremental load must therefore use `UpdatedAt` , not only `CreatedAt` .
- 1 invoice (`ADM06022859336,963953449`) has `create_date` 2025-02-28 but `invoice_date` 2025-03-31 (document created before its date – pre-dated draft). The business date for the fact is `invoice_date` ; flagged in `QualityFlags` , not rejected.
- Invoice keys are `CHAR(25)` with trailing blanks and two prefixes (`ADM…` – 36, `cq0…` – 1) `RTRIM` before lookup.

2.5 Amounts, currency and shipping

- `currency` holds only `'EURO'` (non-ISO) mapped to `EUR` .
- Shipping cost is 2.5 % of `products_total` in 35 invoices; 2 invoices differ (`ADM06022868165,094314602` : 27.70 but `total_for_vat − products_total` = 25.20; `ADM06043059438,135478354` : 26.275 vs 21.275) **rule S1**: when `shipping_cost ≠ total_for_vat − products_total` , use `total_for_vat − products_total` (the header's own arithmetic, which also drives the VAT and final total).
- 2 invoices have `products_total ≠ Σ line_total` (145 vs 140; 234 vs 239). In both the line totals are consistent with `vat_value` (`Σ line_vat = 23 % × Σ line_total`), so **lines are the source of truth** for the fact; the header value is only kept in staging.
- 4 lines have a negative `quantity` or `unit_price` while `line_total` is positive sign error, auto-corrected with `ABS()` .
- All lines: implied VAT = 23 % = `products.vat_rate` .

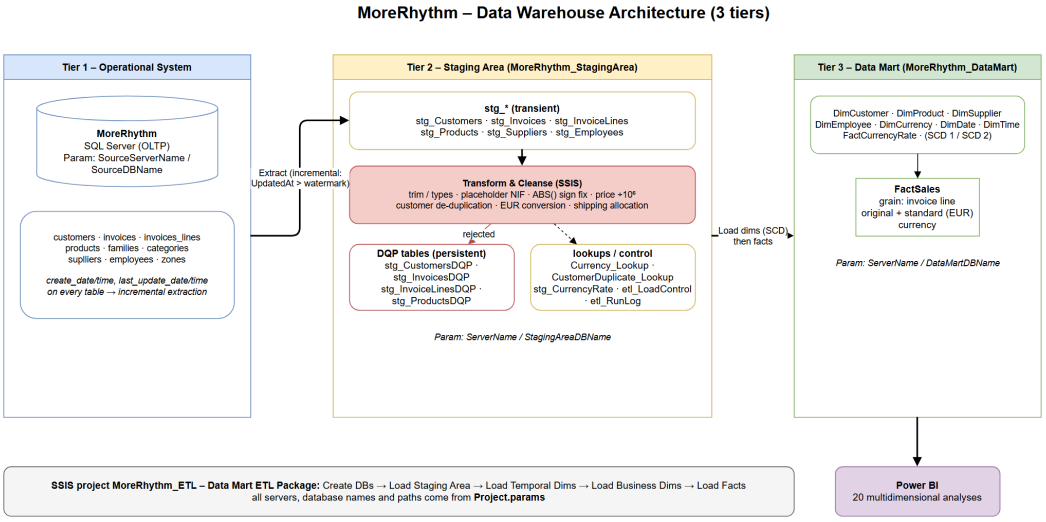
2.6 Products / suppliers

- `current_sell_price` and `last_cost_price` are stored ×1 000 000 (e.g. 680 000 000 vs. line price 680) divided by 10^6.
- 2 "ghost" products (`Z001` , `CORSR`) with every attribute NULL and never sold `stg_ProductsDQP` .
- 10 products without supplier kept with the "Unknown supplier" member (`SupplierKey = -1`).
- Suppliers 8 and 14 share NIF `502199378` ("FaroInstrumentos" / "FaroInstrumentos II", different cities) same company, two warehouses: kept as two suppliers (documented, not merged).

---

3. Data Warehouse Architecture

Three-tier architecture: operational system Staging Area Data Mart, consumed by Power BI (diagram: `Design_Diagrams/Architecture_3Tier.drawio`).



Tier 1 – Operational system (`MoreRhythm`). The SQL Server database that supports billing, purchasing and stock. Every table records `create_date/time` and `last_update_date/time` , which makes incremental extraction possible. Server and database name are parameters (`SourceServerName` , `SourceDBName`).

Tier 2 – Staging Area (`MoreRhythm_StagingArea`). Intermediate database where the data is extracted, cleaned and integrated before it reaches the Data Mart. It holds the transient `stg_*` tables, the persistent `stg_*DQP` tables with the records that cannot be corrected automatically, the `_Lookup` tables, and the ETL control tables (`etl_LoadControl` , `etl_RunLog`).

Tier 3 – Data Mart (`MoreRhythm_DataMart`). Star schema at the invoice-line grain: `FactSales` , `FactCurrencyRate` and the dimensions Customer, Product, Supplier, Employee, Currency, Date and Time. It is the source for the Power BI analyses. The Staging Area and the Data Mart are on the **same server** (`ServerName`), as the brief requires.

ETL flow (SSIS project `MoreRhythm_ETL` , one package `Data Mart ETL Package` organised in sequence containers, as in the labs):

1. **Create databases and tables**: creates the Staging Area and Data Mart databases if they don't exist and runs the `SQL Scripts` for their tables.
2. **Load Staging Area**: extracts only the rows whose `UpdatedAt` is greater than the watermark in `etl_LoadControl` (plus their parent/child rows) into the `stg_*` tables.
3. **Transform & cleanse**: type conversion, trimming, placeholder handling, sign correction, price-scale correction, currency conversion, customer de-duplication and shipping allocation. Records that cannot be corrected automatically go to the `stg_*DQP` table with a `DQP` message and leave the flow (Conditional Split / Lookup error output Derived Column "Add … DQP" Union All OLE DB Destination, as in lab 4).
4. **Load Temporal Dimensions**: `DimDate` and `DimTime` , loaded once from CSV files.
5. **Load Business Dimensions**: Slowly Changing Dimension component (Type 1 / Type 2 per attribute).
6. **Load Fact Tables**: `FactCurrencyRate` , then `FactSales` with surrogate-key lookups. New lines are inserted and changed lines updated (matched on `InvoiceLineID`). The watermarks only move forward if the run succeeds.

4. Dimensional Model

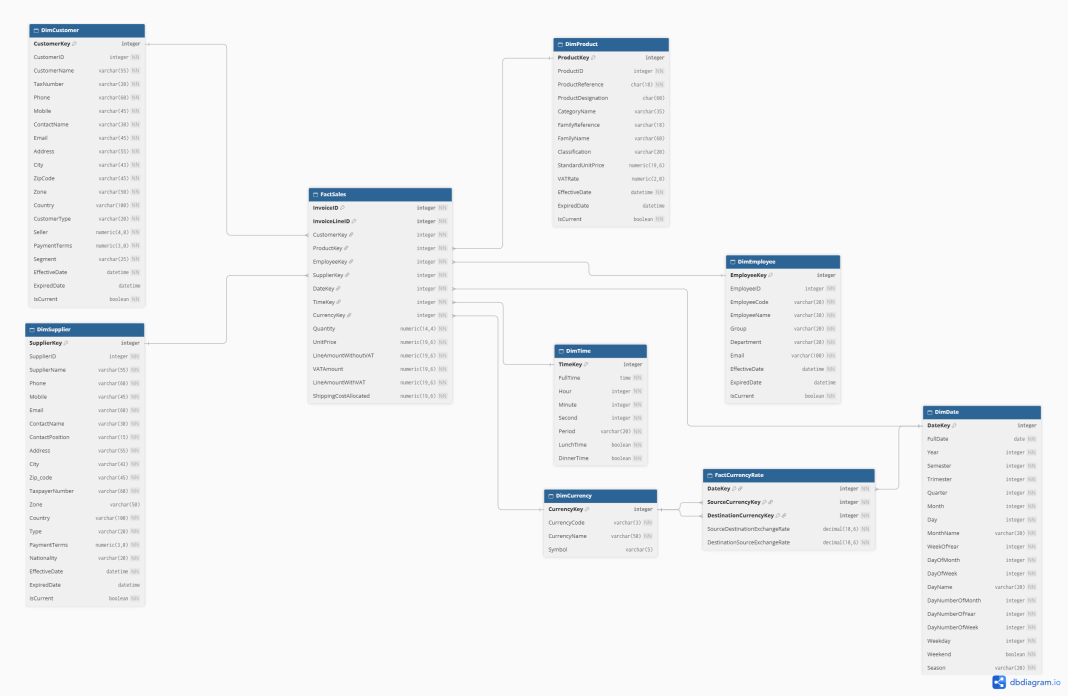
4.1 Kimball methodology

The design of the Data Mart follows the four-step dimensional design process defined by Ralph Kimball:

1. **Business Process:** Sales and billing operations of musical instruments and accessories by More Rhythm, SA.
2. **Granularity Level:** The lowest level of detail chosen for the fact table is **each individual invoice line item** (`invoices_lines`), ensuring maximum analytical flexibility.
3. **Dimensions:** Conformed dimensions providing descriptive context: Customer, Product, Employee, Supplier, Currency, Date, and Time.
4. **Facts / Measures:** Quantitative additive metrics per invoice line (`Quantity` , `UnitPrice` , `LineAmountWithoutVAT` , `VATAmount` , `LineAmountWithVAT` , and `ShippingCostAllocated`).

4.2 Star schema

The conceptual schema is organized as a **Star Schema** centered around the primary fact table `FactSales` , directly connected to the surrounding conformed dimensions via surrogate keys (`CustomerKey` , `ProductKey` , `EmployeeKey` , `SupplierKey` , `DateKey` , `TimeKey` , `CurrencyKey`). An auxiliary fact table `FactCurrencyRate` is included to support historical multi-currency conversions. Snowflake structures are avoided to maximize query performance in Power BI.



4.3 Dimensions

- **`DimCustomer` :** Stores customer details.
- **`DimProduct` :** Contains product attributes, integrating product families and categories hierarchies (e.g., *Ar*, *Cordas*, *Membranas* categories; *Acessórios*, *Amplificadores*, *Teclados* families).
- **`DimEmployee` :** Represents sales representatives and internal staff.
- **`DimSupplier` :** Stores supplier background and geographical zoning.
- **`DimCurrency` :** Supports multi-currency operations and EUR standardisation.
- **`DimDate` & `DimTime` :** Temporal dimensions providing calendar hierarchies and time periods.

4.4 Fact tables

- **`FactSales` :** The core transactional table containing foreign keys to all dimensions, degenerate dimensions (`InvoiceID` , `InvoiceLineID`), and quantitative measures: `Quantity` , `UnitPrice` , `LineAmountWithoutVAT` , `VATRate` , `VATAmount` , `LineAmountWithVAT` , and `ShippingCostAllocated` .
- **`FactCurrencyRate` :** Stores daily exchange rates between source currencies and EUR (`SourceDestinationExchangeRate` , `DestinationSourceExchangeRate`) linked via `DateKey` .

4.5 Attributes and data types

The following tables specify the physical schema of the Data Mart, detailing the attributes, SQL Server data types, nullability constraints, and key assignments (Primary Keys, Foreign Keys, and Degenerate Dimensions) for all dimension and fact tables. Surrogate keys (`*Key`) are implemented as auto-incrementing integers to ensure stability and decoupling from operational keys, while monetary metrics use high-precision `DECIMAL / NUMERIC` types to prevent rounding discrepancies during aggregations.

DimCustomer

Column	Type	Null	Key
CustomerKey	int	no	PK
CustomerID	int	no	
CustomerName	varchar(55)	no	
TaxNumber	varchar(20)	no	
Phone	varchar(60)	no	
Mobile	varchar(45)	no	
ContactName	varchar(30)	no	
Email	varchar(45)	no	
Address	varchar(55)	no	
City	varchar(43)	no	
ZipCode	varchar(45)	no	
Zone	varchar(50)	no	
Country	varchar(100)	no	
CustomerType	varchar(20)	no	
Seller	numeric(4,0)	no	
PaymentTerms	numeric(3,0)	no	
Segment	varchar(25)	no	
EffectiveDate	datetime	no	
ExpiredDate	datetime	yes	
IsCurrent	bit	no	

DimProduct

Column	Type	Null	Key
ProductKey	int	no	PK
ProductReference	char(18)	no	
ProductDesignation	char(60)	yes	
CategoryName	varchar(35)	yes	
FamilyReference	varchar(18)	yes	
FamilyName	varchar(60)	yes	
Classification	varchar(20)	yes	
StandardUnitPrice	numeric(19,6)	yes	

VATRate	numeric(2,0)	yes	
EffectiveDate	datetime	no	
ExpiredDate	datetime	yes	
IsCurrent	bit	no	

DimEmployee

Column	Type	Null	Key
EmployeeKey	int	no	PK
EmployeeID	int	no	
EmployeeCode	varchar(20)	no	
EmployeeName	varchar(30)	no	
Role	varchar(20)	no	
Department	varchar(20)	no	
Email	varchar(100)	no	
EffectiveDate	datetime	no	
ExpiredDate	datetime	yes	
IsCurrent	bit	no	

DimSupplier

Column	Type	Null	Key
SupplierKey	int	no	PK
SupplierID	int	no	
SupplierName	varchar(55)	no	
Phone	varchar(60)	no	
Mobile	varchar(45)	no	
Email	varchar(60)	no	
ContactName	varchar(30)	no	
ContactPosition	varchar(15)	no	
Address	varchar(55)	no	
City	varchar(43)	no	
ZipCode	varchar(45)	no	
TaxpayerNumber	varchar(60)	no	
Zone	varchar(50)	yes	
Country	varchar(100)	no	
Type	varchar(20)	no	
PaymentTerms	numeric(3,0)	no	
Nationality	varchar(20)	no	
EffectiveDate	datetime	no	

ExpiredDate	datetime	yes	
IsCurrent	bit	no	

DimCurrency

Column	Type	Null	Key
CurrencyKey	int	no	PK
CurrencyCode	varchar(3)	no	
CurrencyName	varchar(50)	no	
Symbol	varchar(5)	yes	

DimDate

Column	Type	Null	Key
DateKey	int	no	PK
FullDate	date	no	
Year	int	no	
Semester	int	no	
Quarter	int	no	
Month	int	no	
MonthName	varchar(20)	no	
WeekOfYear	int	no	
DayNumberOfMonth	int	no	
DayNumberOfYear	int	no	
DayNumberOfWeek	int	no	
Weekday	varchar(15)	no	
Weekend	varchar(3)	no	
Season	varchar(20)	no	

DimTime

Column	Type	Null	Key
TimeKey	int	no	PK
FullTime	time	no	
Hour	int	no	
Minute	int	no	
Second	int	no	
Period	varchar(20)	no	
LunchTime	varchar(3)	no	
DinnerTime	varchar(3)	no	

FactSales

Column	Type	Null	Key

InvoiceID	varchar(25)	no	PK / DD
InvoiceLineID	varchar(25)	no	PK / DD
CustomerKey	int	no	FK
ProductKey	int	no	FK
EmployeeKey	int	no	FK
SupplierKey	int	no	FK
DateKey	int	no	FK
TimeKey	int	no	FK
CurrencyKey	int	no	FK
Quantity	numeric(14,4)	no	
UnitPriceOriginalCurrency	numeric(19,6)	no	
LineAmountWithoutVATOriginalCurrency	numeric(19,6)	no	
VATAmountOriginalCurrency	numeric(19,6)	no	
LineAmountWithVATOriginalCurrency	numeric(19,6)	no	
ShippingCostAllocatedOriginalCurrency	numeric(19,6)	no	
UnitPriceStandardCurrency	numeric(19,6)	no	
LineAmountWithoutVATStandardCurrency	numeric(19,6)	no	
VATAmountStandardCurrency	numeric(19,6)	no	
LineAmountWithVATStandardCurrency	numeric(19,6)	no	
ShippingCostAllocatedStandardCurrency	numeric(19,6)	no	

FactCurrencyRate

Column	Type	Null	Key
DateKey	int	no	PK / FK
SourceCurrencyKey	int	no	PK / FK
DestinationCurrencyKey	int	no	PK / FK
SourceDestinationExchangeRate	decimal(18,6)	no	
DestinationSourceExchangeRate	decimal(18,6)	no	

5. Staging Area

Created by `02_Create_Staging_Tables.sql` in database `MoreRhythm_StagingArea`.

Table	Kind	Content
`stg_Customers`	transient	cleaned customers (+ `SurvivorCustomerID`, normalised name, NIF validity)
`stg_Invoices`	transient	invoice headers, ISO currency, rate, corrected shipping, `SumLinesTotal`
`stg_InvoiceLines`	transient	**fact grain**: line measures, allocated shipping, all EUR measures
`stg_Products`, `stg_Suppliers`, `stg_Employees`	transient	dimension sources, already joined with family/category/zone/manager

stg_CustomersDQP , stg_InvoicesDQP , stg_InvoiceLinesDQP , stg_ProductsDQP	persistent	Data Quality Problem tables (course convention, lab 4 CardDQP): rejected rows (raw values as NVARCHAR) + DQP (problem code and description), DQPDate , RunID
CustomerDuplicate_Lookup	persistent	source customer id surviving id (de-duplication result)
Currency_Lookup	persistent	source currency text ISO code (EURO EUR)
stg_CurrencyRate	transient	daily exchange rates, source of FactCurrencyRate
etl_LoadControl	persistent	incremental watermark (LastUpdatedAt) per source table
etl_RunLog	persistent	one row per ETL execution

Every staging/error row carries CreatedAt / UpdatedAt (source *_date + *_time combined into DATETIME2) and RunID for lineage.

5.1 Table structures

The staging tables keep the source business keys (CustomerID , InvoiceID , InvoiceLineID , ProductID , ...) so that the ETL can look up surrogate keys and match changed rows. CreatedAt / UpdatedAt drive the incremental load, and RunID links each row to its row in etl_RunLog .

stg_Customers

Column	Type	Null	Key
CustomerID	numeric(10,0)	no	PK
SurvivorCustomerID	numeric(10,0)	yes	
CustomerName	nvarchar(55)	no	
CustomerNameNorm	nvarchar(55)	yes	
TaxNumber	varchar(20)	yes	
IsTaxNumberValid	bit	yes	
Phone	varchar(20)	yes	
Mobile	varchar(20)	yes	
ContactName	nvarchar(30)	yes	
Email	varchar(45)	yes	
Address	nvarchar(55)	yes	
City	nvarchar(43)	yes	
ZipCode	varchar(10)	yes	
ZipLocality	nvarchar(40)	yes	
Country	nvarchar(30)	yes	
ZoneCode	numeric(5,0)	yes	
ZoneName	nvarchar(50)	yes	
CustomerType	nvarchar(20)	no	
Segment	nvarchar(25)	no	
SellerCode	numeric(4,0)	yes	
DaysToPay	numeric(3,0)	yes	
BBAN	varchar(28)	yes	

CreatedAt	datetime2(0)	no	
UpdatedAt	datetime2(0)	no	
RunID	int	yes	

stg_Invoices

Column	Type	Null	Key
InvoiceID	varchar(25)	no	PK
InvoiceDate	date	no	
CustomerID	numeric(10,0)	no	
SurvivorCustomerID	numeric(10,0)	yes	
EmployeeID	numeric(6,0)	no	
PaymentDate	date	yes	
CurrencySource	varchar(11)	no	
CurrencyCode	char(3)	no	
ExchangeRateToEUR	decimal(19,10)	no	
ProductsTotal	decimal(19,6)	no	
ShippingCost	decimal(19,6)	no	
TotalForVAT	decimal(19,6)	no	
VATValue	decimal(19,6)	no	
FinalTotal	decimal(19,6)	no	
SumLinesTotal	decimal(19,6)	yes	
ShippingCostEUR	decimal(19,6)	yes	
QualityFlags	varchar(200)	yes	
CreatedAt	datetime2(0)	no	
UpdatedAt	datetime2(0)	no	
RunID	int	yes	

stg_InvoiceLines

Column	Type	Null	Key
InvoiceLineID	varchar(25)	no	PK
InvoiceID	varchar(25)	no	
ProductID	varchar(18)	no	
Quantity	decimal(14,4)	no	
UnitPrice	decimal(19,6)	no	
LineTotal	decimal(19,6)	no	
LineVAT	decimal(19,6)	no	
VATRate	decimal(5,2)	yes	
ShippingAllocated	decimal(19,6)	yes	

UnitPriceEUR	decimal(19,6)	yes	
LineTotalEUR	decimal(19,6)	yes	
LineVATEUR	decimal(19,6)	yes	
LineTotalWithVATEUR	decimal(19,6)	yes	
ShippingAllocatedEUR	decimal(19,6)	yes	
QualityFlags	varchar(200)	yes	
CreatedAt	datetime2(0)	no	
UpdatedAt	datetime2(0)	no	
RunID	int	yes	

stg_Products

Column	Type	Null	Key
ProductID	varchar(18)	no	PK
ProductName	nvarchar(60)	no	
FamilyCode	varchar(18)	yes	
FamilyName	nvarchar(60)	yes	
FamilyClassification	nvarchar(20)	yes	
CategoryCode	int	yes	
CategoryName	nvarchar(35)	yes	
SupplierID	numeric(10,0)	yes	
Classification	nvarchar(20)	yes	
CurrentSellPrice	decimal(19,6)	yes	
LastCostPrice	decimal(19,6)	yes	
VATRate	decimal(5,2)	yes	
CreatedAt	datetime2(0)	yes	
UpdatedAt	datetime2(0)	yes	
RunID	int	yes	

stg_Suppliers

Column	Type	Null	Key
SupplierID	numeric(10,0)	no	PK
SupplierName	nvarchar(55)	no	
TaxNumber	varchar(20)	yes	
City	nvarchar(43)	yes	
ZipCode	varchar(10)	yes	
Country	nvarchar(30)	yes	
ZoneCode	numeric(5,0)	yes	
ZoneName	nvarchar(50)	yes	

SupplierType	nvarchar(20)	yes	
Nationality	nvarchar(20)	yes	
CreatedAt	datetime2(0)	yes	
UpdatedAt	datetime2(0)	yes	
RunID	int	yes	

stg_Employees

Column	Type	Null	Key
EmployeeID	numeric(6,0)	no	PK
EmployeeCode	varchar(20)	no	
EmployeeName	nvarchar(30)	no	
EmployeeGroup	nvarchar(20)	yes	
Department	nvarchar(20)	yes	
ManagerID	numeric(6,0)	yes	
ManagerName	nvarchar(30)	yes	
Email	varchar(100)	yes	
CreatedAt	datetime2(0)	yes	
UpdatedAt	datetime2(0)	yes	
RunID	int	yes	

stg_CurrencyRate

Column	Type	Null	Key
CurrencyCode	nchar(3)	no	PK
RateDate	date	no	PK
RateToEUR	decimal(18,6)	no	

Currency_Lookup

Column	Type	Null	Key
SourceCurrency	nvarchar(11)	no	PK
CurrencyCode	nchar(3)	no	

CustomerDuplicate_Lookup

Column	Type	Null	Key
SourceCustomerID	numeric(10,0)	no	PK
SurvivorCustomerID	numeric(10,0)	no	
MatchRule	nvarchar(30)	no	
MatchScore	decimal(5,4)	yes	
RunID	int	yes	

etl_LoadControl

Column	Type	Null	Key
SourceTable	varchar(50)	no	PK
LastUpdatedAt	datetime2(0)	no	
LastRunID	int	yes	
LastRunTime	datetime2(0)	yes	

etl_RunLog

Column	Type	Null	Key
RunID	int	no	PK
PackageName	nvarchar(128)	no	
StartTime	datetime2(0)	no	
EndTime	datetime2(0)	yes	
Status	varchar(10)	no	
RowsExtracted	int	yes	
RowsRejected	int	yes	
RowsLoaded	int	yes	

DQP tables (`stg_CustomersDQP` , `stg_InvoicesDQP` , `stg_InvoiceLinesDQP` , `stg_ProductsDQP`) follow the lab convention (`CardDQP`). Each has the same columns as the source table, stored as `NVARCHAR` so that values with the wrong type or length can still be kept, plus:

Column	Type	Null	Description
DQPID	int identity	no	primary key
DQP	nvarchar(200)	no	problem code and description, e.g. `C01: customer name is missing`
DQPDate	datetime	no	when the record was rejected (default `GETDATE()`)
RunID	int	yes	ETL execution that rejected it

These tables are never truncated. Correcting the records and loading them later is out of the scope of this work, as the brief states.

6. Data Mapping and SCD Strategy

6.1 Mapping: source Staging Area Data Mart

The logical mapping below traces every attribute from the operational source (`MoreRhythm`), through the Staging Area (`MoreRhythm_StagingArea`), into the target Data Mart (`MoreRhythm_DataMart`), specifying all transformation rules, formulas, and associated Data Quality (DQ) codes defined in Section 7.1.

*(Note: Temporal dimensions `DimDate` and `DimTime` are **generated** once from pre-built CSV files rather than mapped from operational tables, while surrogate keys and SCD Type 2 metadata columns— `EffectiveDate` , `ExpiredDate` , `IsCurrent` —are generated during ETL execution.)*

---

1. `DimCustomer` Mapping

Source (MoreRhythm)	Staging (MoreRhythm_StagingArea)	Data Mart (MoreRhythm_DataMart)	Transformation / Cleansing Rule	DQ / Tra Cod

(Generated)	*(Generated)*	DimCustomer.CustomerKey	Auto-increment Surrogate Key (SK)	—
customers.number	stg_Customers.SurvivorCustomerID	DimCustomer.CustomerID	De-duplication clustering; survivor ID selected via rules M1/M2	C05, C06
customers.name	stg_Customers.CustomerName	DimCustomer.CustomerName	LTRIM/RTRIM; reject if NULL/empty	C01, C11
customers.taxpayer_number	stg_Customers.TaxNumber	DimCustomer.TaxNumber	LTRIM/RTRIM; placeholders ('', 0, 000000000) NULL in staging 'N/A' in DW; modulo-11 check flagged (IsTaxNumberValid)	C02, C07, C11
customers.phone	stg_Customers.Phone	DimCustomer.Phone	LTRIM/RTRIM; empty 'N/A'; survivorship recovery	C11
customers.mobile	stg_Customers.Mobile	DimCustomer.Mobile	LTRIM/RTRIM; empty 'N/A'; survivorship recovery	C11
customers.contact	stg_Customers.ContactName	DimCustomer.ContactName	LTRIM/RTRIM; empty 'N/A'; survivorship recovery	C11
customers.email	stg_Customers.Email	DimCustomer.Email	LTRIM/RTRIM; invalid format (no @) NULL recovered from duplicate or 'N/A'	C03, C11
customers.address	stg_Customers.Address	DimCustomer.Address	LTRIM/RTRIM; empty 'N/A'	C11
customers.clity	stg_Customers.City	DimCustomer.City	LTRIM/RTRIM; if missing, parsed from zip_code (Spain) or set to 'Unknown'	C08, C11
customers.zip_code	stg_Customers.ZipCode	DimCustomer.ZipCode	Split on first space after NNNN-NNN (keeps code part); foreign codes kept as-is	C09, T10
customers.zone + zones.designation	stg_Customers.ZoneName	DimCustomer.Zone	Lookup on zones; missing 'Unknown'	C04, T01
(Derived from Zone)	stg_Customers.Country	DimCustomer.Country	Derived: ZoneName = 'Espanha' 'Spain', other valid zones 'Portugal', missing 'Unknown'	C04, T02

customers.type	stg_Customers.CustomerType	DimCustomer.CustomerType	LTRIM/RTRIM ; empty 'N/A'	C11
customers.seller	stg_Customers.SellerCode	DimCustomer.Seller	Kept as descriptive attribute only (not a FK; 0 flagged)	C10
customers.days_to_pay	stg_Customers.DaysToPay	DimCustomer.PaymentTerms	Direct type cast to NUMERIC(3,0)	—
customers.segment	stg_Customers.Segment	DimCustomer.Segment	LTRIM/RTRIM ; empty 'N/A'	C11
(ETL Metadata)	*(ETL Metadata)*	DimCustomer.EffectiveDate	GETDATE() (or 1900-01-01 on initial load)	—
(ETL Metadata)	*(ETL Metadata)*	DimCustomer.ExpiredDate	NULL while active; updated to GETDATE() when superseded	—
(ETL Metadata)	*(ETL Metadata)*	DimCustomer.IsCurrent	1 for active version, 0 for historical versions	—

---

2. DimProduct Mapping

Source (MoreRhythm)	Staging (MoreRhythm_StagingArea)	Data Mart (MoreRhythm_DataMart)	Transformation / Cleansing Rule	DQ / Transf. Code
(Generated)	*(Generated)*	DimProduct.ProductKey	Auto-increment Surrogate Key (SK)	—
products.reference	stg_Products.ProductID	DimProduct.ProductID	RTRIM trailing blanks from CHAR(18)	P04
products.designation	stg_Products.ProductName	DimProduct.ProductDesignation	LTRIM/RTRIM ; reject row if all descriptive fields are NULL	P01, C11
products.category + categories.designation	stg_Products.CategoryName	DimProduct.CategoryName	Lookup on categories ; empty 'Unknown'	T04
products.family	stg_Products.FamilyCode	DimProduct.FamilyReference	RTRIM ; empty 'N/A'	P04
products.family + families.designation	stg_Products.FamilyName	DimProduct.FamilyName	Lookup on families ; empty 'Unknown'	T03
families.classification	stg_Products.FamilyClassification	DimProduct.Classification	Lookup on families ; empty 'N/A'	T03
products.current_sell_price	stg_Products.CurrentSellPrice	DimProduct.StandardUnitPrice	Scale correction: current_sell_price / 1000000	P02
products.vat_rate	stg_Products.VATRate	DimProduct.VATRate	Direct numeric cast	L05

			(NUMERIC(2,0))	
(ETL Metadata)	*(ETL Metadata)*	DimProduct.EffectiveDate	GETDATE() (or 1900-01-01 on initial load)	—
(ETL Metadata)	*(ETL Metadata)*	DimProduct.ExpiredDate	NULL while active; updated to GETDATE() when superseded	—
(ETL Metadata)	*(ETL Metadata)*	DimProduct.IsCurrent	1 for active version, 0 for historical versions	—

---

3. DimEmployee Mapping

Source (MoreRhythm)	Staging (MoreRhythm_StagingArea)	Data Mart (MoreRhythm_DataMart)	Transformation / Cleansing Rule	DQ / Transf. Code
(Generated)	*(Generated)*	DimEmployee.EmployeeKey	Auto-increment Surrogate Key (SK)	—
employees.number	stg_Employees.EmployeeID	DimEmployee.EmployeeID	Direct numeric cast (INTEGER)	—
employees.code	stg_Employees.EmployeeCode	DimEmployee.EmployeeCode	LTRIM/RTRIM	C11
employees.name	stg_Employees.EmployeeName	DimEmployee.EmployeeName	LTRIM/RTRIM	C11
employees.[group]	stg_Employees.EmployeeGroup	DimEmployee.Role	LTRIM/RTRIM ; mapped to role description	T06, C11
employees.department	stg_Employees.Department	DimEmployee.Department	LTRIM/RTRIM ; empty 'N/A'	C11
employees.email	stg_Employees.Email	DimEmployee.Email	LTRIM/RTRIM ; empty 'N/A'	C11
(ETL Metadata)	*(ETL Metadata)*	DimEmployee.EffectiveDate	GETDATE() (or 1900-01-01 on initial load)	—
(ETL Metadata)	*(ETL Metadata)*	DimEmployee.ExpiredDate	NULL while active; updated to GETDATE() when superseded	—
(ETL Metadata)	*(ETL Metadata)*	DimEmployee.IsCurrent	1 for active version, 0 for historical versions	—

---

4. DimSupplier Mapping

Source (MoreRhythm)	Staging (MoreRhythm_StagingArea)	Data Mart (MoreRhythm_DataMart)	Transformation / Cleansing Rule	DQ / Transf. Code
(Generated)	*(Generated)*	DimSupplier.SupplierKey	Auto-increment SK (-1 reserved	P03

			for 'Unknown Supplier')	
supplliers.number	stg_Suppliers.SupplierID	DimSupplier.SupplierID	Direct numeric cast (INTEGER)	—
supplliers.name	stg_Suppliers.SupplierName	DimSupplier.SupplierName	LTRIM/RTRIM	C11
supplliers.phone	stg_Suppliers.Phone *(raw)*	DimSupplier.Phone	LTRIM/RTRIM; empty 'N/A'	C11
supplliers.mobile	stg_Suppliers.Mobile *(raw)*	DimSupplier.Mobile	LTRIM/RTRIM; empty 'N/A'	C11
supplliers.email	stg_Suppliers.Email *(raw)*	DimSupplier.Email	LTRIM/RTRIM; empty 'N/A'	C11
supplliers.contact	stg_Suppliers.ContactName *(raw)*	DimSupplier.ContactName	LTRIM/RTRIM; empty 'N/A'	C11
supplliers.contact_position	stg_Suppliers.ContactPosition *(raw)*	DimSupplier.ContactPosition	LTRIM/RTRIM; empty 'N/A'	C11
supplliers.address	stg_Suppliers.Address *(raw)*	DimSupplier.Address	LTRIM/RTRIM; empty 'N/A'	C11
supplliers.city	stg_Suppliers.City	DimSupplier.City	LTRIM/RTRIM; missing 'Unknown'	S02, C11
supplliers.zip_code	stg_Suppliers.ZipCode	DimSupplier.Zip_code	LTRIM/RTRIM; split code/locality where applicable	T10, C11
supplliers.taxpayer_number	stg_Suppliers.TaxNumber	DimSupplier.TaxpayerNumber	LTRIM/RTRIM; shared NIFs kept separate if different cities	S01, C11
supplliers.zone + zones.designation	stg_Suppliers.ZoneName	DimSupplier.Zone	Lookup on zones; missing 'Unknown'	S02, T01
(Derived from Zone)	stg_Suppliers.Country	DimSupplier.Country	Derived: 'Espanha' 'Spain', other valid zones 'Portugal', missing 'Unknown'	S02, T02
supplliers.type	stg_Suppliers.SupplierType	DimSupplier.Type	LTRIM/RTRIM; empty 'N/A'	C11
supplliers.days_to_pay	stg_Suppliers.DaysToPay *(raw)*	DimSupplier.PaymentTerms	Direct cast to NUMERIC(3,0)	—
supplliers.nationality	stg_Suppliers.Nationality	DimSupplier.Nationality	LTRIM/RTRIM; empty 'N/A'	C11
(ETL Metadata)	*(ETL Metadata)*	DimSupplier.EffectiveDate	GETDATE() (or 1900-01-01 on initial load)	—
(ETL Metadata)	*(ETL Metadata)*	DimSupplier.ExpiredDate	NULL while active; updated	—

			to GETDATE() when superseded	
(ETL Metadata)	*(ETL Metadata)*	DimSupplier.IsCurrent	1 for active version, 0 for historical versions	—

5. DimCurrency Mapping

Source (MoreRhythm / Lookup)	Staging (MoreRhythm_StagingArea)	Data Mart (MoreRhythm_DataMart)	Transformation / Cleansing Rule	DQ / Transf. Code
(Generated)	*(Generated)*	DimCurrency.CurrencyKey	Auto-increment Surrogate Key (SK)	—
invoices.currency	Currency_Lookup.CurrencyCode	DimCurrency.CurrencyCode	Map free text ('EURO' 'EUR' , 'GBP' 'GBP'); reject unknown	I06, T07
(Reference Lookup)	Currency_Lookup	DimCurrency.CurrencyName	Derived from ISO code ('EUR' 'Euro' , 'GBP' 'Pound Sterling')	T07
(Reference Lookup)	Currency_Lookup	DimCurrency.Symbol	Derived from ISO code ('EUR' '€' , 'GBP' '£')	T07

6. DimDate & DimTime (Generated Dimensions)

Target Table	Generation Method & Derived Attributes	DQ / Transf. Code
DimDate	**Generated** once at initial deployment from Files\DimDate.csv (SQL Server Analysis Services time table generation). Not mapped from operational tables. Includes derived attributes: DateKey (YYYYMMDD), FullDate , Year , Semester , Trimester , Quarter , Month , Day , MonthName , WeekOfYear , DayOfMonth , DayOfWeek , DayName , DayNumberOfMonth , DayNumberOfYear , DayNumberOfWeek , Weekday , Weekend ('Yes' / 'No'), and Season (Spring , Summer , Autumn , Winter based on meteorological/astronomical date ranges).	T08
DimTime	**Generated** once at initial deployment from Files\DimTime.csv (86,400 daily seconds). Not mapped from operational tables. Includes derived attributes: TimeKey (HHMMSS), FullTime , Hour , Minute , Second , Period (Morning , Afternoon , Evening , Night), LunchTime ('Yes' / 'No' for 12:00–14:30), and DinnerTime ('Yes' / 'No' for 19:30–22:00).	T08

7. FactSales Mapping

Source (MoreRhythm)	Staging (MoreRhythm_StagingArea)	Data Mart (MoreRhythm_DataMart)	Transformati / Formula Rul
invoices.invoice_stamp_id	stg_Invoices.InvoiceID	FactSales.InvoiceID	RTRIM trailin Degenerate Di

invoices_lines.invoice_line_stamp_id	stg_InvoiceLines.InvoiceLineID	FactSales.InvoiceLineID	RTRIM trailin Degenerate Di Grain)
invoices.customer	stg_Invoices.SurvivorCustomerID	FactSales.CustomerKey	Resolved via CustomerDup then SCD Type DimCustome InvoiceDate
invoices_lines.product	stg_InvoiceLines.ProductID	FactSales.ProductKey	RTRIM ; SCD on DimProdu InvoiceDate ; missing/P01
invoices.employee	stg_Invoices.EmployeeID	FactSales.EmployeeKey	SCD Type 2 loo DimEmployee InvoiceDate ; found
products.supplyer	stg_Products.SupplierID	FactSales.SupplierKey	SCD Type 2 loo DimSupplier ProductID ; d (Unknown) if
invoices.invoice_date	stg_Invoices.InvoiceDate	FactSales.DateKey	Formatted as i YYYYMMDD MONTH*100 + match DimDa reject if NULL/
invoices.create_time	stg_Invoices.CreatedAt *(time part)*	FactSales.TimeKey	Formatted as i HHMMSS (H MINUTE*100 + match DimTir
invoices.currency	stg_Invoices.CurrencyCode	FactSales.CurrencyKey	Lookup on DimCurrency ('EURO' 'E
invoices_lines.quantity	stg_InvoiceLines.Quantity	FactSales.Quantity	Sign correctio ABS(quantity) negative with line_total
invoices_lines.unit_price	stg_InvoiceLines.UnitPriceEUR	FactSales.UnitPrice	ABS(unit_pric ExchangeRate
invoices_lines.line_total	stg_InvoiceLines.LineTotalEUR	FactSales.LineAmountWithoutVAT	line_total * ExchangeRate validating `
invoices_lines.line_vat	stg_InvoiceLines.LineVATEUR	FactSales.VATAmount	If NULL: (line_ VATRate) * ExchangeRate line_vat * ExchangeRate
(Derived)	stg_InvoiceLines.LineTotalWithVATEUR	FactSales.LineAmountWithVAT	Formula: (Lin LineVAT) * ExchangeRate
invoices.shipping_cost	stg_InvoiceLines.ShippingAllocatedEUR	FactSales.ShippingCostAllocated	Proportional a rounding resic adjustment: F

			(LineTotal / SU 6) * ExchangeI (where Ship* ProductsTotal shipping_cost

---

8. `FactCurrencyRate` Mapping

Source (External ECB File)	Staging (MoreRhythm_StagingArea)	Data Mart (MoreRhythm_DataMart)	Transformation / Cle / Formula Rule
(Generated)	*(Generated)*	FactCurrencyRate.CurrencyRateKey	Auto-increment Surrog Key (BIGINT)
CurrencyRates.csv (RateDate)	stg_CurrencyRate.RateDate	FactCurrencyRate.DateKey	Formatted as integer YYYYMMDD to matcl DimDate.DateKey
CurrencyRates.csv (SourceCode)	stg_CurrencyRate.CurrencyCode	FactCurrencyRate.SourceCurrencyKey	Lookup against DimCurrency.Currenc
CurrencyRates.csv (TargetCode)	*(Fixed 'EUR' / Target)*	FactCurrencyRate.DestinationCurrencyKey	Lookup against DimCurrency.Currenc
CurrencyRates.csv (RateToEUR)	stg_CurrencyRate.RateToEUR	FactCurrencyRate.SourceDestinationExchangeRate	Direct cast to DECIMAL(18,6) (1.00 for EUR -> EUR)
(Derived)	*(Derived)*	FactCurrencyRate.DestinationSourceExchangeRate	Inverse rate formula: ROUND(1.0 / NULLIF(RateToEUR, 0),

6.2 SCD Strategy (Slowly Changing Dimensions)

To handle changes in dimensional attributes over time, the SSIS Slowly Changing Dimension component configuration is structured using **Type 1 (Overwriting)** and **Type 2 (Historical Versioning)** strategies.

Summary of SCD Attribute Classification per Dimension

1. `DimCustomer`

- **Fixed / Business Key:** `CustomerID` (Unchanging natural identifier).
- **Changing (Type 1 - Overwriting):** `Phone` , `Mobile` , `ContactName` , `Email` .
  - *Justification:* Contact details are operational attributes; correcting a typo or updating a phone number should overwrite current records without altering past sales context.
- **Historical (Type 2 - Versioning):** `CustomerName` , `TaxNumber` , `Address` , `City` , `ZipCode` , `Zone` , `Country` , `CustomerType` , `Seller` , `PaymentTerms` , `Segment` .
  - *Justification:* These attributes define geographical grouping, customer segmentation, and business behavior. If a customer changes address or moves to another sales zone, historical sales must remain associated with the zone and address active at the exact time of the transaction.

2. `DimProduct`

- **Fixed / Business Key:** `ProductID` .
- **Changing (Type 1 - Overwriting):** `StandardUnitPrice` .
  - *Justification:* Standard prices can be overwritten for future transactions, while historical transaction pricing is safely captured by the unit price measures directly inside the fact table (`FactSales`).
- **Historical (Type 2 - Versioning):** `ProductDesignation` , `CategoryName` , `FamilyReference` , `FamilyName` , `Classification` , `VATRate` .
  - *Justification:* Products moving between categories or families must preserve historical reporting integrity. Filtering sales by family/category across different periods requires Type 2 versioning.

3. `DimEmployee`

- **Fixed / Business Key:** `EmployeeID` , `EmployeeCode` .
- **Changing (Type 1 - Overwriting):** `EmployeeName` , `Email` .
  - *Justification:* Minor corrections to a name spelling or email updates do not impact historical sales performance evaluations.
- **Historical (Type 2 - Versioning):** `Role` , `Department` .
  - *Justification:* Employees transferring between departments or organizational roles must have their sales performance correctly attributed to the department they belonged to at the time of each sale.

4. **`DimSupplier`**

- **Fixed / Business Key:** `SupplierID` .
- **Changing (Type 1 - Overwriting):** `Phone` , `Mobile` , `Email` , `ContactName` , `ContactPosition` .
  - *Justification:* Purely operational supplier contact data updated via Type 1.
- **Historical (Type 2 - Versioning):** `SupplierName` , `Address` , `City` , `Zip_code` , `TaxpayerNumber` , `Zone` , `Country` , `Type` , `PaymentTerms` , `Nationality` .
  - *Justification:* Supplier location and contractual classification affect geographic and supply-chain reporting over time.

5. **`DimCurrency`** , **`DimDate`** , **&** **`DimTime`**

- **Fixed (Type 0 / Static):** All attributes in `DimCurrency` , `DimDate` , and `DimTime` are treated as static/fixed. Dates, times, and ISO currency definitions represent immutable temporal or international standards and do not change.

Justification for Excluding Types 3, 4, and 6

- **Type 3 (Previous Value columns):** Not used because it only preserves a single prior version (e.g., "Previous Address"), which is inadequate for multi-year historical analyses where customers might change locations multiple times.
- **Type 4 (Mini-Dimension / History Table):** Not used because splitting rapidly changing attributes into a separate mini-dimension introduces unnecessary complexity and joins, whereas Type 2 row-versioning handles the historical depth cleanly within the main dimension tables.
- **Type 6 (Hybrid):** Not used because combining Type 1, 2, and 3 behaviours introduces excessive complexity and redundant data maintenance, whereas strict separation between Type 1 (operational edits) and Type 2 (analytical history) satisfies all business requirements cleanly.

Technical Execution of Type 2 Versioning

When a Type 2 attribute changes:

1. The existing active record is expired by updating `IsCurrent = 0` and setting `ExpiredDate = GETDATE()` .
2. A new record is inserted with the updated attribute values, `IsCurrent = 1` , `EffectiveDate = GETDATE()` , and a newly generated surrogate key (`CustomerKey` , `ProductKey` , etc.).
3. **Fact Lookup Mechanism:** When loading `FactSales` , the foreign key lookup matches the source business entity and evaluates the transaction date (`DateKey` mapped to `FullDate`), ensuring the fact table links exclusively to the dimension version valid on the exact invoice date (`FullDate BETWEEN EffectiveDate AND ISNULL(ExpiredDate, '9999-12-31')`).

7. Data Quality, Cleansing and Transformations

7.1 Data quality problems and how each is handled

Every problem found by profiling (or that the ETL must guard against on future data) gets one of three treatments:

- **AUTO-CORRECT**: the ETL can repair the value deterministically. The record is loaded. For invoices and lines the fix is also noted in the `QualityFlags` column of `stg_Invoices` / `stg_InvoiceLines` .
- **FLAG**: the value is suspicious but not wrong enough to block the record. The record is loaded unchanged and flagged.
- **REJECT**: no automatic correction is possible. The record is written to the matching `stg_<Table>DQP` table with a `DQP` message (problem code + description), removed from `stg_*` , and **never loaded into the Data Mart**, as the brief requires. Correcting and reloading rejected records is out of scope.

"Rows in sample" is the number of records affected in the data provided (`01_Data_Profiling_results.txt`); 0 means the rule is a safeguard for future loads.

Customers (`customers` → `stg_Customers` / `stg_CustomersDQP`)

Code	Problem	Attribute(s)	Rows in sample	Treatment	How
C01	Name missing	name	0	REJECT `stg_CustomersDQP`	a customer cannot be identified without a name
C02	Placeholder tax number (`''`, `0`, `00`, `000000000`)	taxpayer_number	4 (ids 2, 20, 24, 25)	AUTO-CORRECT	`TaxNumber = NULL` in staging, `'N/A'` in DimCustomer; record kept (internal branches / cash sales); excluded from NIF matching
C03	Invalid e-mail format (no `@`)	email	1 (id 5: `coro.sapo.pt`)	AUTO-CORRECT	`Email = NULL`; then recovered from duplicate id 26 (`coro@sapo.pt`) by survivorship
C04	Zone missing	zone	2 (ids 2, 20)	AUTO-CORRECT	`Zone = 'Unknown'`, `Country = 'Unknown'`
C05	Exact duplicate (same NIF)	taxpayer_number	1 group (ids 10, 18)	AUTO-CORRECT (merge)	de-duplication rule M1 (Section 7.2)
C06	Near-duplicate (NIF typo, name variant, same zip)	name, taxpayer_number, zip_code	2 groups (5 26, 8 22)	AUTO-CORRECT (merge)	de-duplication rule M2 (Section 7.2)
C07	NIF fails modulo-11 check digit	taxpayer_number	17 of 23	FLAG	`IsTaxNumberValid = 0`; used only to pick the survivor (rejecting would discard most customers)
C08	City missing	clity	3 (ids 2, 20, 21)	AUTO-CORRECT	id 21: city parsed from zip text ("Alcobendas - Madrid"); others `'Unknown'`
C09	Zip code and locality stored together; foreign format	zip_code	24 / 1 foreign (id 21)	AUTO-CORRECT	split into `ZipCode` (`NNNN-NNN`) + `ZipLocality`; foreign codes kept as they are
C10	Seller code does not match an employee	seller	4 (seller = 0)	FLAG	kept as a descriptive attribute only, not used as a foreign key
C11	Text with extra spaces / inconsistent case	all text	all	AUTO-CORRECT	`LTRIM/RTRIM`; empty string NULL in staging, `'N/A'` in the Data Mart

Products and suppliers (`products`, `suplliers` → `stg_Products`, `stg_Suppliers` / `stg_ProductsDQP`)

Code	Problem	Attribute(s)	Rows in sample	Treatment	How
P01	Product with every descriptive attribute NULL	designation, family, category	2 (`Z001`, `CORSR`)	**REJECT `stg_ProductsDQP`**	never sold, so nothing cascades to sales

P02	Price stored ×1 000 000	current_sell_price, last_cost_price	65 (all non-null)	AUTO-CORRECT	÷ 1 000 000 (checked against line `unit_price`: ratio = 10^6)
P03	Product without supplier	supplyer	10	AUTO-CORRECT	`SupplierKey = -1` (Unknown member of DimSupplier)
P04	Product reference with trailing blanks (`CHAR(18)`)	reference	all	AUTO-CORRECT	`RTRIM` before lookups
S01	Two suppliers share a NIF	taxpayer_number	1 pair (8, 14 "FaroInstrumentos"/"II")	FLAG	different cities, so treated as two warehouses of one company; not merged
S02	Supplier zone / city missing	zone, city	2 (3 "Compras a Dinheiro", 13 internal)	AUTO-CORRECT	`Zone`, `City`, `Country = 'Unknown'`

Invoices (`invoices` `stg_Invoices` / `stg_InvoicesDQP`)

Code	Problem	Attribute(s)	Rows in sample	Treatment	How
I01	Invoice date NULL or in the future	invoice_date	0	REJECT `stg_InvoicesDQP`	no valid date means no `DateKey`
I02	Invoice without lines	—	0	REJECT `stg_InvoicesDQP`	nothing to load at the fact grain
I03	Customer or employee not found	customer, employee	0	REJECT `stg_InvoicesDQP`	broken foreign key
I04	Shipping ≠ `total_for_vat − products_total`	shipping_cost	2	AUTO-CORRECT	rule S1: use `total_for_vat − products_total`
I05	`products_total` ≠ Σ `line_total`	products_total	2	FLAG	lines are consistent with the VAT, so lines are used for the fact
I06	Currency not ISO / unknown / no exchange rate	currency	37 non-ISO (`EURO`), 0 unknown	AUTO-CORRECT (`EURO` `EUR`); unknown currency REJECT	mapping through `Currency_Lookup`
I07	Record created before its invoice date	create_date	1	FLAG	`invoice_date` is the business date for the fact
I08	Payment date before invoice date	payment_date	0	FLAG	safeguard
I09	Created after updated	create_, *last_update_*	0	FLAG	safeguard (would break the incremental watermark)
I10	Invoice key padded to `CHAR(25)`	invoice_stamp_id	37	AUTO-CORRECT	`RTRIM`

Invoice lines (`invoices_lines` `stg_InvoiceLines` / `stg_InvoiceLinesDQP`)

Code	Problem	Attribute(s)	Rows in sample	Treatment	How
L01	Negative quantity or unit price while `line_total` is positive	quantity, unit_price	4	AUTO-CORRECT	`ABS()` (sign entry error; `line_total` and VAT confirm the positive value)
L02	Product not found or rejected (P01)	product	0	REJECT `stg_InvoiceLinesDQP`	no `ProductKey`
L03	`	quantity × unit_price	≠ line_total` (beyond 0.01)	line_total	0
L04	Line belongs to a rejected invoice	invoice_stamp_id	0	REJECT `stg_InvoiceLinesDQP` (cascade)	keeps fact and header consistent
L05	VAT NULL or ≠ `line_total × vat_rate`	line_vat	0 (all 23 %)	NULL AUTO-CORRECT (recompute from rate); mismatch FLAG	

Result on the provided sample: 2 records are rejected (products `Z001` and `CORSR`, which were never sold). Every customer, invoice and invoice line can be corrected automatically, so the Data Mart receives 23 customers (after merging 3 duplicate pairs), 37 invoices and all 93 lines, with shipping and amounts reconciled.

Transformations that make values explicit (no data quality problem)

As in lab 4 (`Gender` F Female), every code is replaced by a value whose meaning is clear at analysis time:

Code	Source	Data Mart	Transformation
T01	`customers.zone`, `suplliers.zone` (1–5)	`Zone`	Lookup on `zones` Norte, Centro, Sul, Ilhas, Espanha
T02	zone	`Country`	derived: Espanha Spain, other zones Portugal
T03	`products.family` (reference)	`FamilyName`, `FamilyClassification`	Lookup on `families`
T04	`products.category` (code)	`CategoryName`	Lookup on `categories`
T05	`employees.chief` (number)	`ManagerName`	self-lookup on `employees`; `'N/A'` for top management
T06	`employees.[group]`	`Role`	renamed attribute (Vendedor, Administrativo, Administrador)
T07	`invoices.currency` (`EURO`)	`DimCurrency.Code`, `Name`, `Symbol`	`Currency_Lookup` ISO code DimCurrency
T08	dates / times	`Weekend`, `LastDayOfMonth`, `LunchTime`, `DinnerTime`	`'Yes'` / `'No'` instead of 1/0
T09	`create_date` + `create_time`, `last_update_date` + `last_update_time`	`CreatedAt`, `UpdatedAt` (staging)	combined into one `DATETIME2`
T10	`zip_code` (`2600-011 VILA FRANCA DE XIRA`)	`PostalCode` + locality	split on the first space after the `NNNN-NNN` code

7.2 Customer de-duplication

Input: `stg_Customers` (new/changed customers) and the customers already in `DimCustomer`. **Output:** `CustomerDuplicate_Lookup` (every source id → survivor id) and one surviving, merged record per real customer.

1. **Standardise** every record:
   - `NameNorm` = UPPER, accent-insensitive (`COLLATE Latin1_General_CI_AI`), punctuation removed, legal suffixes removed (`LDA`, `S.A.`, `UNIPESSOAL`), abbreviations expanded (`N.` → `NOSSA`, `SRA` → `SENHORA`, `R.` → `RUA`, `AV.` → `AVENIDA`), multiple spaces collapsed;
   - `AddressNorm` the same way; `Phone` / `Mobile` → digits only;
   - `ZipCode` = first 8 characters (`NNNN-NNN`);
   - `TaxNumber` = NULL if blank or only zeros (rule C02); `IsTaxNumberValid` = modulo-11 check.
2. **Exclude non-matchable records** from the comparison: customers of type `Interno` and generic customers ("Vendas a Dinheiro") – each is its own survivor.
3. **Blocking** (to avoid comparing all n^2 pairs on large volumes): only compare records that share **NIF**, or **ZipCode**, or **Phone**.
4. **Matching rules** for each candidate pair (A, B):
   - **M1 – exact identifier:** `A.TaxNumber = B.TaxNumber` (not NULL) → duplicate, score 1.0.
   - **M2 – fuzzy:** same `ZipCode` **and** name similarity ≥ 0.80 (SSIS *Fuzzy Grouping* on `NameNorm`, `AddressNorm`) **and** at least one corroborating evidence: NIF differs in ≤ 2 digits, **or** same phone, **or** address similarity ≥ 0.80 → duplicate, score = similarity.
   - Otherwise → distinct customers.
5. **Clustering:** take the transitive closure of the matched pairs (if A≈B and B≈C then {A,B,C} is one group).
6. **Survivor selection** inside each group, in order:
   1. record already present in `DimCustomer` (keeps the surrogate key stable);
   2. record with a valid NIF check digit;
   3. most complete record (largest number of non-empty attributes);
   4. earliest `CreatedAt` (the original registration);
   5. lowest `CustomerID`.
7. **Survivorship per attribute:** survivor values win; empty/invalid survivor attributes are filled from the other members with the most recent `UpdatedAt` (e.g. customer 5's invalid e-mail is replaced by customer 26's `coro@sapo.pt`).
8. **Persist** `CustomerDuplicate_Lookup(SourceCustomerID, SurvivorCustomerID, MatchRule, MatchScore)`; only survivors are loaded into `DimCustomer`.
9. **Re-key facts:** `stg_Invoices.SurvivorCustomerID` is set through the map, so sales of all duplicates are attributed to the same `CustomerKey`.
10. **Incremental runs:** new customers are matched against `DimCustomer` as well; an existing map entry is never re-pointed automatically (manual review).

Result on the provided data (validated in SQL): 3 groups, 0 false positives – {10, 18} by M1, {5, 26} and {8, 22} by M2. 26 source customers → 23 customers in the Data Mart.

Survivor of each group (rules of step 6 applied to the data; none of them is in `DimCustomer` yet, and none has a valid check digit except group 1, where both NIFs are equal and valid):

Group	Members	Rule	Completeness (filled contact/address fields)	Survivor	Reason
1	10, 18 "Eira e Beira, lda"	M1	6 / 6	**10**	tie on completeness → created first (2025-03-31 vs 2025-05-31)
2	5 "Coro Nossa Senhora Assunção", 26 "Coro N. Senhora Assunção"	M2	5 / 6	**26**	more complete (valid e-mail, mobile); customer 5's contact name is kept by attribute survivorship
3	8 "Luz e Som, lda", 22 "Luz e Som Lda."	M2	5 / 5	**8**	tie on completeness → created first (2025-04-30 vs 2025-06-30); phone taken from 22

Validation query (finds the candidate pairs with rules M1 and M2; result: exactly the 3 pairs above):

```
WITH c AS (
 SELECT number, zip_code,
       LTRIM(RTRIM(REPLACE(REPLACE(REPLACE(REPLACE(REPLACE(REPLACE(
          UPPER(name) COLLATE Latin1_General_CI_AI, ',', ' '), '.', ' '), ' LDA', ' '),
          ' S A', ' '), ' N ', ' NOSSA '), '  ', ' '))) AS NameNorm,
       taxpayer_number AS nif,
       CASE WHEN taxpayer_number LIKE '%[^0]%' AND LEN(taxpayer_number) = 9 THEN 1 ELSE 0 END AS real_nif
 FROM MoreRhythm.dbo.customers
)
SELECT a.number AS customer_a, b.number AS customer_b,
     CASE WHEN a.real_nif = 1 AND a.nif = b.nif THEN 'M1' ELSE 'M2' END AS match_rule
FROM c a JOIN c b ON a.number < b.number
WHERE (a.real_nif = 1 AND a.nif = b.nif)
  OR (a.zip_code = b.zip_code AND a.zip_code <> '' AND DIFFERENCE(a.NameNorm, b.NameNorm) = 4);
```

In SSIS, the same matching is done with the **Fuzzy Grouping** transformation on `NameNorm`, `AddressNorm` and `ZipCode` (exact match on `ZipCode`), with a similarity threshold of 0.80.

7.3 Shipping-cost allocation

Shipping is recorded only on the invoice header; the fact grain is the invoice line, so the cost is allocated **proportionally to the line value excluding VAT**.

For invoice *i* with lines *j = 1…n* :

$$\text{Ship}_i^* = \begin{cases} \text{TotalForVAT}_i - \text{ProductsTotal}_i & \text{if ShippingCost}_i \models \text{TotalForVAT}_i - \text{ProductsTotal}_i \text{ (rule S1)} \\ \text{ShippingCost}_i & \text{otherwise} \end{cases}$$

$$\text{ShippingAllocated}_{ij} = \text{round}_6\left(\text{Ship}_i^* \times \frac{\text{LineTotal}_{ij}}{\sum_{k=1}^{n_i} \text{LineTotal}_{ik}}\right)$$

Rounding residual (so the line shares add up exactly to the header):

$$\text{ShippingAllocated}_{i,j^*} \mathrel{+}= \text{Ship}_i^* - \sum_{j=1}^{n_i} \text{ShippingAllocated}_{ij}, \qquad j^* = \arg\max_j \text{LineTotal}_{ij}$$

$$\text{ShippingCostOriginalCurrency}_{ij} = \text{ShippingAllocated}_{ij}, \qquad \text{ShippingCostStandardCurrency}_{ij} = \text{ShippingAllo}$$

Why proportional and not an equal split. The slides (“Header & Line Facts with Different Granularities”) use the lab example to show that the header value must not be copied to every line. Copying freight 15 to 3 lines would sum to 45. The example splits it equally (5 + 5 + 5). We use the same allocation idea, but weight by line value. A 1 322 € guitar line and a 56 € case line on the same invoice should not carry the same shipping cost. Both methods keep the total exact (Σ lines = header).

Notes

- `LineTotal` is used **after** the sign correction (L01), so weights are never negative.
- If `Σ LineTotal = 0` (all lines free of charge) the cost is split by quantity: `Ship* × Quantity_ij / Σ Quantity_ik`.
- The denominator is the **sum of the lines** (not the header `products_total`, which is wrong in 2 invoices), guaranteeing `Σ_j ShippingAllocated_ij = Ship*_i`.
- Worked example – invoice `ADM06013159066,541488606`, shipping 51.45, lines 680 + 56 + 1322 = 2058: 680/2058 × 51.45 = 17.00; 56/2058 × 51.45 = 1.40; 1322/2058 × 51.45 = 33.05; total 51.45.
- Validation on the full sample (37 invoices, 93 lines): Σ allocated = 990.395 = Σ invoice shipping, and the line shares of every invoice add up exactly to its shipping cost (0 mismatches).

SQL implementation (Transform & Cleanse step of the `Data Mart ETL Package`, Execute SQL Task on the Staging Area):

```
WITH l AS (
 SELECT l.InvoiceLineID, l.InvoiceID, l.LineTotal, i.ShippingCost,
       SUM(l.LineTotal) OVER (PARTITION BY l.InvoiceID) AS SumLines,
       ROW_NUMBER() OVER (PARTITION BY l.InvoiceID ORDER BY l.LineTotal DESC, l.InvoiceLineID) AS rn
```

```
  FROM stg_InvoiceLines l JOIN stg_Invoices i ON i.InvoiceID = l.InvoiceID
), a AS (
  SELECT *, ROUND(ShippingCost * LineTotal / NULLIF(SumLines,0), 6) AS Alloc FROM l
), b AS (
  SELECT InvoiceLineID,
         Alloc + CASE WHEN rn = 1 THEN ShippingCost - SUM(Alloc) OVER (PARTITION BY InvoiceID) ELSE 0 END AS AllocFinal
  FROM a
)
UPDATE t SET ShippingAllocated = b.AllocFinal
FROM stg_InvoiceLines t JOIN b ON b.InvoiceLineID = t.InvoiceLineID;
```

8. Currency Standardisation and SSIS Parameterisation

8.1 EUR multi-currency standardisation

Requirement: the company will expand to other EU countries and to the UK, so the Data Mart uses the **euro as the reporting currency** while still allowing analysis by original currency.

Design (the "Multiple Currencies" technique from the Dimensional Data Modeling slides, also used in the resolution of modelling exercise 4):

1. **Normalise the currency code.** Source values are free text (`'EURO'`). A Lookup against `Currency_Lookup` maps them to ISO-4217 codes (`EURO` → `EUR` ; later `LIBRA` / `GBP` → `GBP` , etc.). A currency with no match cannot be corrected automatically → `stg_InvoicesDQP` (`I06: unknown currency`).
2. **Currency dimension** `DimCurrency(CurrencyKey, Code, Name, Symbol)` : the fact references the currency in which the invoice was issued.
3. **Pairs of facts.** Each money fact of `FactSales` is stored twice: in the **original (local) currency** and in the **standard currency (EUR)**: `UnitPrice…` , `Amount…` (excl. VAT), `VATAmount…` , `TotalAmount…` (incl. VAT), `ShippingCost…` (allocated), each with the suffix `OriginalCurrency` or `StandardCurrency` . EUR facts can be added across all invoices; original facts can be analysed by filtering on one currency.
4. **Currency conversion fact table** `FactCurrencyRate(DateKey, SourceCurrencyKey, DestinationCurrencyKey, SourceDestinationExchangeRate, DestinationSourceExchangeRate)` keeps the daily rates in **both directions**, because the two directions are not always symmetric. This lets analyses be shown in any currency the user selects, not only EUR. Rates are loaded from a daily rates file (e.g. ECB reference rates) into `stg_CurrencyRate` ; EUR → EUR has rate 1 for every date.
5. **Conversion at load time**: the rate used is the one for the invoice date (or the latest earlier date with a rate):

```
SELECT TOP (1) r.SourceDestinationExchangeRate
FROM FactCurrencyRate r
JOIN DimDate d ON d.DateKey = r.DateKey
WHERE r.SourceCurrencyKey = @InvoiceCurrencyKey
  AND r.DestinationCurrencyKey = @EURKey
  AND d.FullDate <= @InvoiceDate
ORDER BY d.FullDate DESC;
-- <Fact>StandardCurrency = <Fact>OriginalCurrency × rate
```

6. The rate is fixed at load time (the historical rate of the invoice date). If a rate is corrected later, the invoice's `UpdatedAt` does not change, so the affected period has to be reloaded manually. This is a documented limitation.

With the current data all invoices are EUR (rate = 1), so both values of every pair are equal. The mechanism is in place for the expansion to the UK.

8.2 SSIS project parameterisation

Requirement: **no absolute paths or server/database names inside components**.
The first four parameters are the ones defined in lab exercise 2 (same names).
The two `Source...` parameters are added because the project brief also requires the operational database server and name to be configurable.

Parameter	Type	Default (dev)	Used by
`ServerName`	String	`localhost`	Staging Area & Data Mart connections (same server)
`StagingAreaDBName`	String	`MoreRhythm_StagingArea`	`CM_StagingArea` ; "Create Staging Area Database" task (`@DBName`)
`DataMartDBName`	String	`MoreRhythm_DataMart`	`CM_DataMart` ; "Create Data Mart Database" task (`@DBName`)
`PackagePath`	String	project folder (where the solution is)	Foreach loops over `SQL Scripts\Staging Area` , `SQL Scripts\Data Mart` , CSV files in `Files`
`SourceServerName`	String	`localhost`	`CM_MoreRhythm` connection
`SourceDBName`	String	`MoreRhythm`	`CM_MoreRhythm` connection

How it is applied (following lab tutorials 1 and 2)

- **Creating the databases**: an Execute SQL Task on an **ADO.NET** connection to `master` runs a dynamic `IF DB_ID(@DBName) IS NULL EXEC('CREATE DATABASE ' + @DBName)` , with `@DBName` mapped to `$Project::StagingAreaDBName` (or `DataMartDBName`).
- **Creating the tables**: a Foreach File Enumerator iterates `@[$Project::PackagePath] + "SQL Scripts\\Staging Area"` (`*.sql`), and an Execute SQL Task runs each script (`CreateStageTable<Name>.sql` ; Data Mart: `CreateDim<Name>.sql` , `CreateFact<Name>.sql`). Every script checks that its table does not already exist first.
- **Temporal dimensions** are loaded once, only if empty, from `@[$Project::PackagePath] + "Files\\DimDate.csv"` (generated with an Analysis Services project, lab 3) and `"Files\\DimTime.csv"` (generated in Excel, lab 2).
- The **project-level OLE DB connection managers** have their `ConnectionString` set by a property **expression**, e.g.

```
"Data Source=" + @[$Project::ServerName] + ";Initial Catalog=" + @[$Project::StagingAreaDBName]
+ ";Provider=MSOLEDBSQL.1;Integrated Security=SSPI;Auto Translate=False;"
```

- No component contains a literal path. Every file or folder is built as `@[$Project::PackagePath]` plus a path relative to the project (`SQL Scripts\...` , `Files\...`).
- All the ETL is in one package (`Data Mart ETL Package`) organised in sequence containers, as in the labs, so no package calls another package by path.
- On another machine (the other team member's or the teacher's), only the values in `Project.params` change. After deploying to the SSIS catalog (SSISDB), the parameters can be mapped to an **SSIS Environment** (DEV/TEST/PROD) without touching the package, as in the "Deploying Packages" tutorial.

9. Incremental Loading

To avoid full reloads on every ETL execution, the pipeline implements an **incremental loading mechanism** combining source audit timestamps (`CreatedAt` and `UpdatedAt`) with business key matching:

1. **Change Detection at Extraction (Source → Staging):**
    - The control table `etl_LoadControl` stores the timestamp of the last successful extraction (`LastUpdatedAt`) for each source table.
    - During extraction, the ETL queries the operational database (`MoreRhythm`) and extracts only records where `UpdatedAt > LastUpdatedAt` . This single filter captures both newly inserted records (since `CreatedAt = UpdatedAt` on creation) and existing records that were recently modified.
2. **New vs. Modified Identification (Staging → Data Mart):**
   Once cleansed in the Staging Area, records are compared against the target Data Mart tables using their natural

business identifiers (e.g., `CustomerID`, `ProductID`, `InvoiceLineID`):

- **New Records (`INSERT`):** If the business ID does not yet exist in the target dimension or fact table, the record is treated as new and inserted with a newly generated Surrogate Key.
- **Modified Records (`UPDATE` / SCD Versioning):** If the business ID already exists in the Data Mart, the record is recognized as an update. Depending on the target table and attribute rules, the ETL either overwrites the values (SCD Type 1 / fact line corrections) or closes the previous version and inserts a new active record (SCD Type 2).

3. **Watermark Update:**
   - Only after all packages complete without errors is `etl_LoadControl.LastUpdatedAt` updated to the maximum timestamp of the processed batch, ensuring no records are lost if an execution fails.

10. Conclusions and Next Steps

Part 1 defined the complete design of the MoreRhythm data mart:

- a three-tier architecture (operational system, Staging Area, Data Mart) loaded by one parameterised SSIS project;
- a star schema at the invoice-line grain, with seven dimensions and two fact tables, that answers the 10 required analyses and any other analysis at that level of detail;
- the Staging Area, with a Data Quality Problem table for every source table whose records can be rejected;
- the handling of every data quality problem found in profiling: 33 rules, of which only the 2 empty products are rejected in the sample;
- the customer de-duplication (26 23 customers), the shipping allocation to the line level, the euro standardisation and incremental loading based on the source timestamps.

In **Part 2** the design will be implemented: the SSIS packages (Staging Area load, cleansing with DQP tables, temporal dimensions, SCD component for the business dimensions, fact loading), the Analysis Services project, and the 20 Power BI analyses.

References

- R. Kimball, M. Ross, *The Data Warehouse Toolkit: The Definitive Guide to Dimensional Modeling*, Wiley.
- R. Kimball, M. Ross, *Relentlessly Practical Tools for Data Warehousing and Business Intelligence*, Wiley, 2015.
- P. Oliveira, *Armazéns de Dados* – lecture slides (Introduction to Data Warehouses; Dimensional Data Modeling), ISEP, 2026/2027.
- ARMDD laboratory tutorials 1–5 (Integration Services), ISEP, 2026/2027.